


W-Shape: Aligned vs Base Model
(Qwen3-0.6B vs Qwen3-0.6B-Base, n=50 per category)



Pairwise Trajectories: Hate Speech Outlier
(Aligned=solid, Base=dashed)

